

# Wine Quality Classification

```
In [1]: import numpy as np
import seaborn as sns
from ucimlrepo import fetch_ucirepo
from sklearn.cluster import KMeans
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
```

```
In [2]: wine_quality = fetch_ucirepo(id=186)
X = wine_quality.data.features
y = wine_quality.data.targets
```

```
In [3]: print(X)
```

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	
...	...	...	...	...	...	...
6492	6.2	0.21	0.29	1.6	0.039	
6493	6.6	0.32	0.36	8.0	0.047	
6494	6.5	0.24	0.19	1.2	0.041	
6495	5.5	0.29	0.30	1.1	0.022	
6496	6.0	0.21	0.38	0.8	0.020	

	free_sulfur_dioxide	total_sulfur_dioxide	density	pH	sulphates	\
0	11.0	34.0	0.99780	3.51	0.56	
1	25.0	67.0	0.99680	3.20	0.68	
2	15.0	54.0	0.99700	3.26	0.65	
3	17.0	60.0	0.99800	3.16	0.58	
4	11.0	34.0	0.99780	3.51	0.56	
...	...	...	...	...	...	...
6492	24.0	92.0	0.99114	3.27	0.50	
6493	57.0	168.0	0.99490	3.15	0.46	
6494	30.0	111.0	0.99254	2.99	0.46	
6495	20.0	110.0	0.98869	3.34	0.38	
6496	22.0	98.0	0.98941	3.26	0.32	

	alcohol
0	9.4
1	9.8
2	9.8
3	9.8
4	9.4
...	...
6492	11.2
6493	9.6
6494	9.4
6495	12.8
6496	11.8

[6497 rows x 11 columns]

```
In [4]: print(y)
```

	quality
0	5
1	5
2	5
3	6
4	5
...	...
6492	6
6493	5
6494	6
6495	7
6496	6

[6497 rows x 1 columns]

## Pre-processing of data

```
In [5]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [6]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

y_train = y_train.values.flatten()
y_test = y_test.values.flatten()
```

## Unsupervised K-Means Clustering Model

```
In [7]: kmeans = KMeans()
kmeans_labels_train = kmeans.fit_predict(X_train_scaled)
```

```
In [8]: cluster_to_label = {}
for cluster in np.unique(kmeans_labels_train):
    cluster_indices = np.where(kmeans_labels_train == cluster)[0] # Find the indices of the current cluster
    most_common_label = np.bincount(y_train[cluster_indices]).argmax() # Get the most common label
    cluster_to_label[cluster] = most_common_label # Label each cluster

kmeans_labels_train_mapped = np.array([cluster_to_label[label] for label in kmeans_labels_train])
```

## Unsupervised K-Means Clustering Model

```
In [9]: kmeans_labels_test = kmeans.predict(X_test_scaled)
kmeans_labels_test_mapped = np.array([cluster_to_label[label] for label in kmeans_labels_test])
```

```
In [10]: correct_predictions = np.sum(kmeans_labels_test_mapped == y_test)
total_predictions = len(y_test)

model_accuracy = correct_predictions / total_predictions
print(f"Accuracy: {model_accuracy:.4f}")
```

Accuracy: 0.4746

```
In [11]: report = classification_report(y_test, kmeans_labels_test_mapped, target_names=[str(i) for i in np.unique(y_test)])
print("Classification Report (K-Means Clustering):\n", report)
```


```
Classification Report (K-Means Clustering):
      precision    recall  f1-score   support

     3         0.00      0.00      0.00         6
     4         0.00      0.00      0.00        43
     5         0.46      0.43      0.44       402
     6         0.48      0.75      0.59       597
     7         0.00      0.00      0.00       215
     8         0.00      0.00      0.00        36
     9         0.00      0.00      0.00         1

 accuracy          0.47          1300
 macro avg         0.13          0.15          1300
 weighted avg      0.36          0.41          1300
```

```
c:\Users\randa\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\_classification.py:1565: Undefined
MetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` param
eter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\randa\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\_classification.py:1565: Undefined
MetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` param
eter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\randa\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\_classification.py:1565: Undefined
MetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` param
eter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
In [12]: confusion_matrix1 = confusion_matrix(y_test, kmeans_labels_test_mapped)
plt.figure(figsize=(10, 7))
sns.heatmap(confusion_matrix1, annot=True, fmt='d', cmap='Blues', xticklabels=np.unique(y_test), yticklabels=np.unique(y_t
plt.title("Confusion Matrix (K-Means Clustering)")
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.show()
```



### Supervised Logistic Regression Classification Model

```
In [13]: log_reg_model = LogisticRegression()
log_reg_model.fit(X_train_scaled, y_train)
```

```
Out[13]: LogisticRegression
LogisticRegression()
```

```
In [14]: y_pred = log_reg_model.predict(X_test_scaled)
```

```
In [15]: correct_predictions = np.sum(y_pred == y_test)
total_predictions = len(y_test)

model_accuracy = correct_predictions / total_predictions
print(f"Logistic Regression Accuracy: {model_accuracy:.4f}")
```

Logistic Regression Accuracy: 0.5362

### Results

```
In [16]: report = classification_report(y_test, y_pred, target_names=[str(i) for i in np.unique(y_test)])
print("Classification Report (Logistic Regression):\n", report)
```

```
Classification Report (Logistic Regression):
              precision    recall  f1-score   support

     3         1.00      0.17      0.29         6
     4         0.00      0.00      0.00        43
     5         0.54      0.61      0.57       402
     6         0.54      0.68      0.60       597
     7         0.50      0.20      0.29       215
     8         0.00      0.00      0.00         36
     9         0.00      0.00      0.00          1

 accuracy          0.54
 macro avg         0.37      0.24      0.25       1300
 weighted avg         0.50      0.54      0.50       1300
```

```
c:\Users\randa\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\_classification.py:1565: Undefined
MetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` param
eter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\randa\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\_classification.py:1565: Undefined
MetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` param
eter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\randa\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\_classification.py:1565: Undefined
MetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` param
eter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
In [17]: confusion_matrix2 = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(10, 7))
sns.heatmap(confusion_matrix2, annot=True, fmt='d', cmap='Blues', xticklabels=np.unique(y_test), yticklabels=np.unique(y_t
plt.title("Confusion Matrix (Logistic Regression)")
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.show()
```

